

# AIFA: Stochastic Planning MDP [Policy Iteration]

09/04/2024

**Koustav Rudra**

# Policy Iteration


- Alternative approach for optimal values:
  - **Step 1: Policy Evaluation:**
    - calculate utilities for some fixed policy (not optimal utilities!) until convergence
  - **Step 2: Policy Improvement:**
    - update policy using one-step look-ahead with resulting converged (but not optimal!) utilities as future values
    - Repeat steps until policy converges

# Policy Evaluation

- How do we calculate the V's for a fixed policy  $\pi$ ?
- Idea 1: Turn recursive Bellman equations into updates

$$V_0^\pi(s) = 0$$

$$V_{k+1}^\pi(s) \leftarrow \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V_k^\pi(s')]$$



# Policy Iteration

- Evaluation: For fixed current policy  $\pi$ , find values with policy evaluation:
  - Iterate until values converge:

$$V_{k+1}^{\pi_i}(s) \leftarrow \sum_{s'} T(s, \pi_i(s), s') \left[ R(s, \pi_i(s), s') + \gamma V_k^{\pi_i}(s') \right]$$

- Improvement: For fixed values, get a better policy using policy extraction
  - One-step look-ahead:

$$\pi_{i+1}(s) = \arg \max_a \sum_{s'} T(s, a, s') \left[ R(s, a, s') + \gamma V^{\pi_i}(s') \right]$$

# Comparison

- Both value iteration and policy iteration compute the same thing (all optimal values)
- In value iteration:
  - Every iteration updates both the values and (implicitly) the policy
  - We don't track the policy, but taking the max over actions implicitly recomputes it
- In policy iteration:
  - We do several passes that update utilities with fixed policy (each pass is fast because we consider only one action, not all of them)
  - After the policy is evaluated, a new policy is chosen (slow like a value iteration pass)
  - The new policy will be better (or we're done)
- Both are dynamic programs for solving MDPs

# Introduction to Learning

09/04/2024

**Koustav Rudra**

# Paradigms of Learning

- **Supervised Learning**

- Both inputs and outputs are given
- The outputs are provided by experts


- **Reinforcement Learning**

- The agent receives some evaluation of actions
- E.g., fine for giving a wrong chess move

- **Unsupervised Learning**


- No Label, no feedback
- Have to learn pattern from the inputs

# Supervised Learning: Classification





# Supervised Learning: Regression



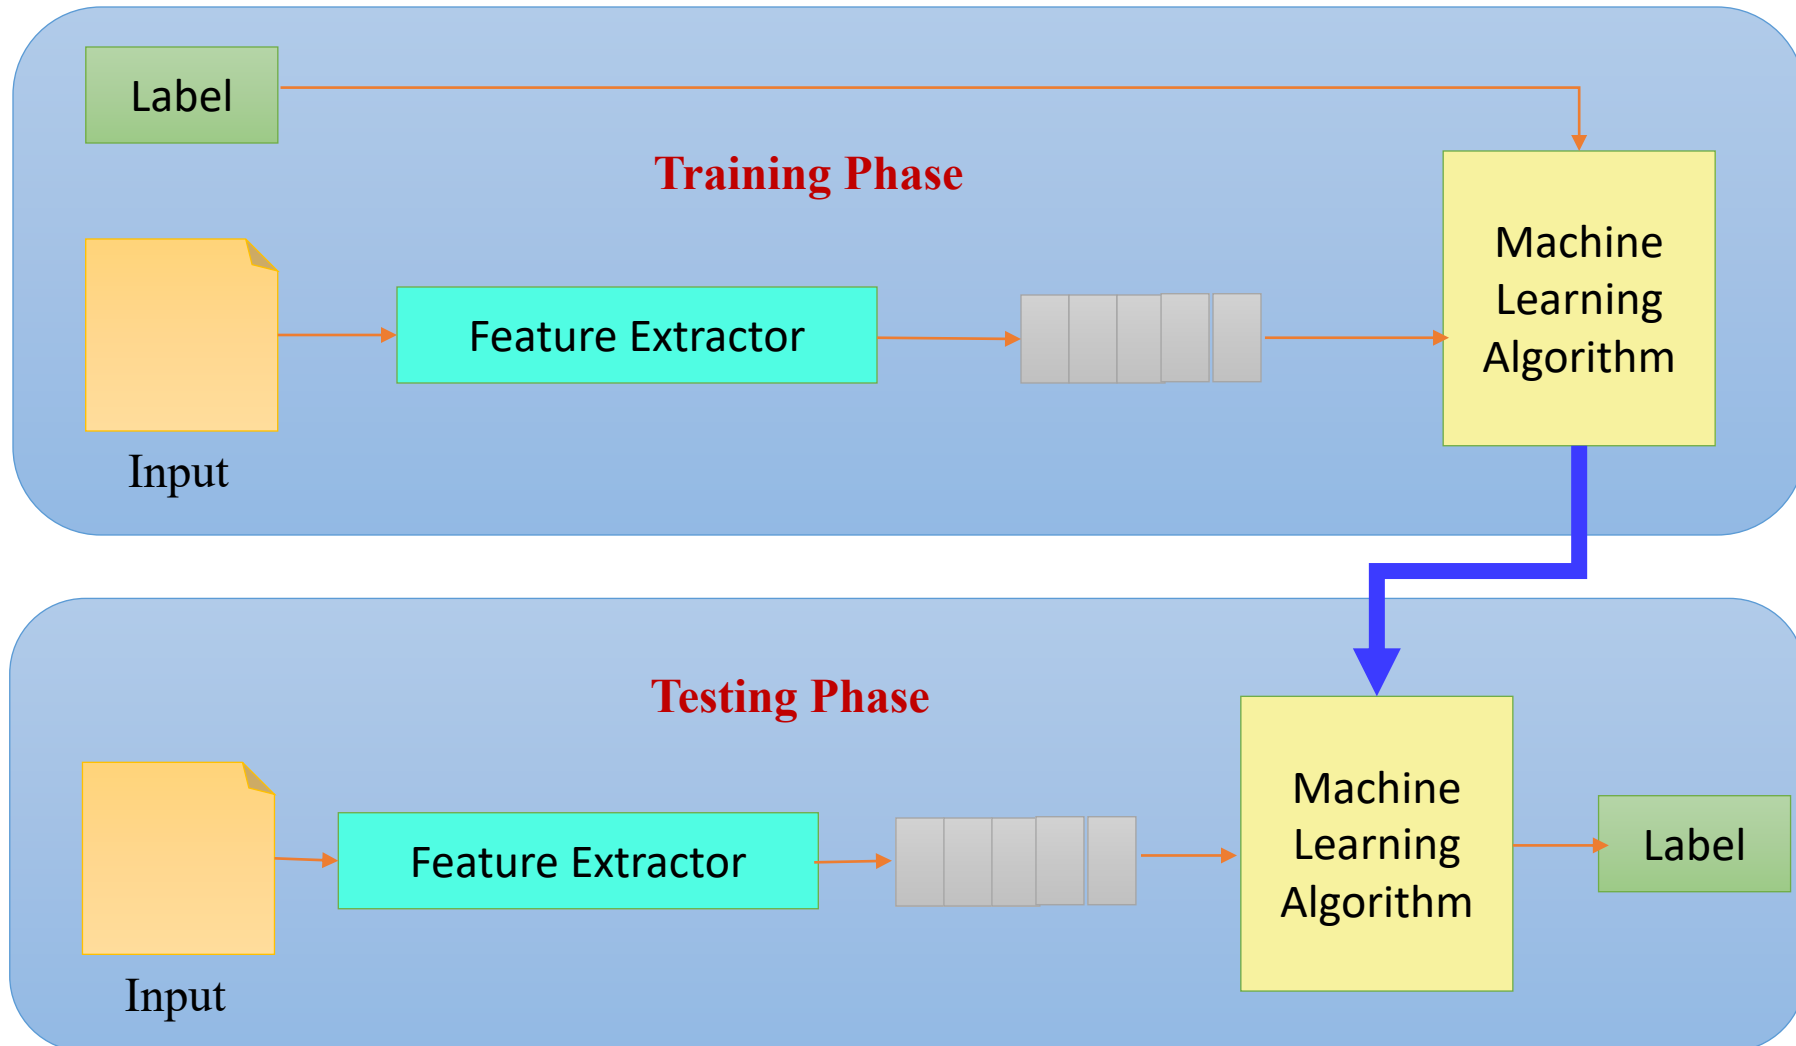
Supervised Learning  
“right answers” given

Regression: Predict  
continuous valued output  
(price)


# Features

- Often the **individual observations** are analyzed into a set of quantifiable properties which are called features
  - **categorical - well-defined finite set of values**
    - nominal (e.g. A, B, AB or O, for blood types)
    - ordinal (e.g. large, medium, small)
    - dichotomous (e.g. male, female)
  - **integer-valued** (e.g. the number of words in a text)
  - **real-valued** (e.g. height)


# Supervised Learning: Classification



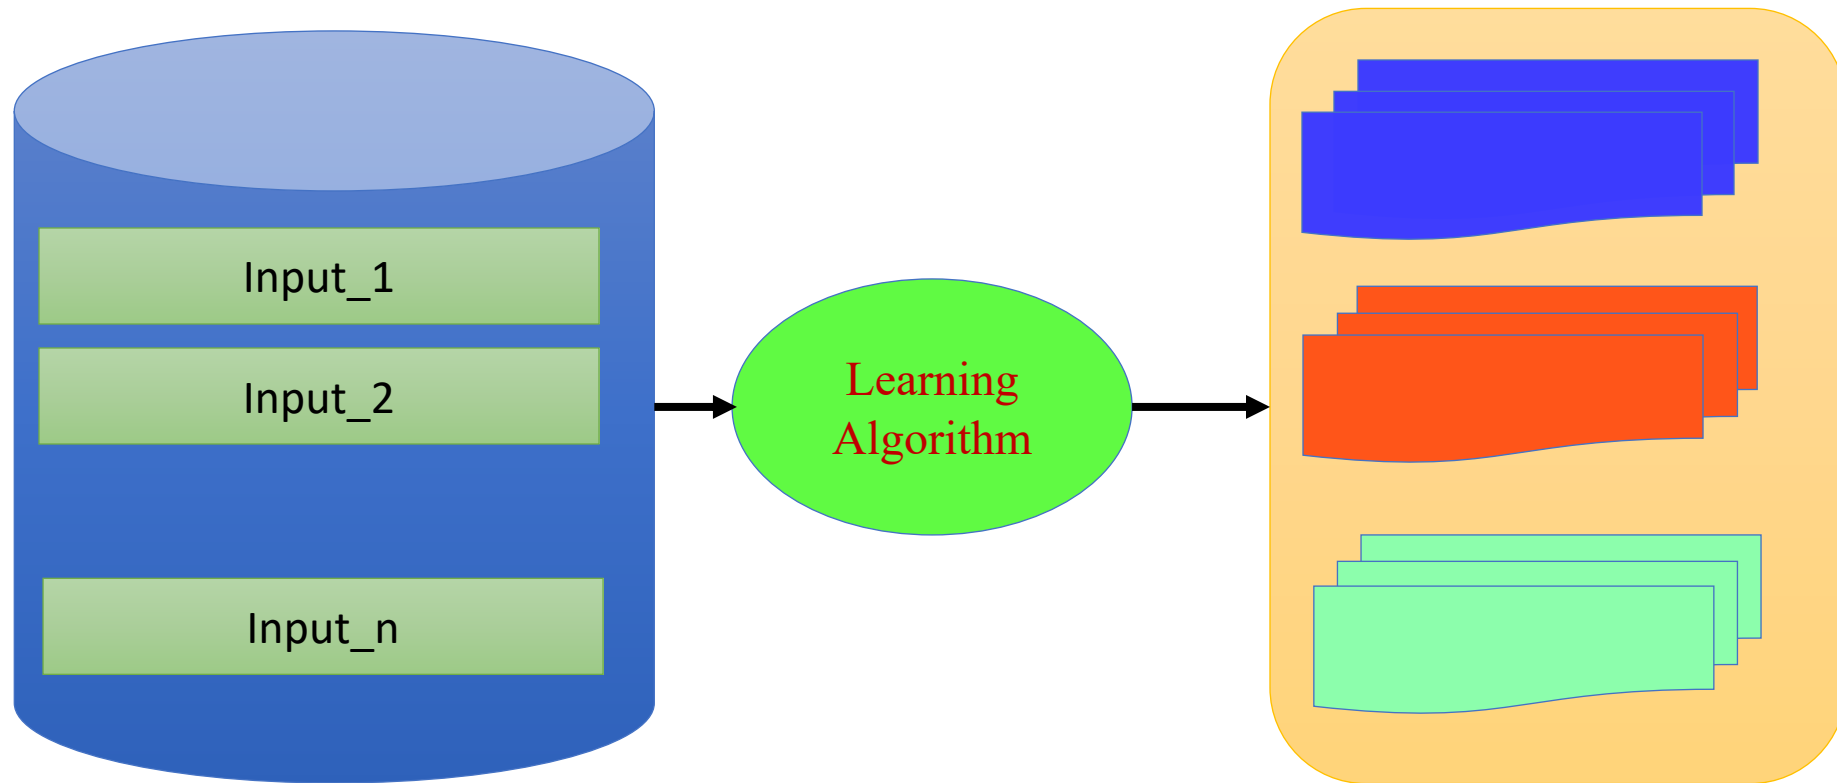
# Unsupervised Learning: Clustering



# Unsupervised Learning: Clustering



# Unsupervised Learning



# Terminology

- **Features:** The number of distinct traits that can be used to describe each item in a quantitative manner
- **Feature vector:** n-dimensional vector of numerical features that represent some object
- **Instance Space  $X$ :** Set of all possible objects describable by features
- **Example  $(x,y)$ :** Instance  $x$  with label  $y=f(x)$

# Concept Learning



# Concept Learning Task

- **Target concept:** “days on which Jack should enjoy sports”

	Sky	AirTemp	Humidity	Wind	Water	EnjoySport
1	Sunny	Warm	Normal	Strong	Warm	Yes
2	Sunny	Warm	High	Strong	Warm	Yes
3	Rainy	Cold	High	Strong	Warm	No
4	Sunny	Warm	High	Strong	Cool	Yes

# Concept Learning Task: Example

- $X = \{1, 2, 3, 4\}$ 
  - Set of all possible days
- **Features:**  $\langle \text{Sky}, \text{AirTemp}, \text{Humidity}, \text{Wind}, \text{Water} \rangle$
- Target concept (c) is to be learned
  - $c(x)=1$  if EnjoySport = Yes
  - $c(x)=0$  if EnjoySport = No
  - $c : \text{EnjoySport}: X \rightarrow \{0, 1\}$
- **Training data D:**  $\{(1,1), (2,1), (3,0), (4,1)\}$

# Hypothesis Space and Inductive Learning

# Concept Learning

- Given examples of a data point  $D = \{(x, c(x))\}$
- Find out a hypothesis  $h: X \rightarrow \{0, 1\}$ 
  - $h$  basically approximates  $c$
- Hypothesis Space:  $H = \{h_1, h_2, h_3, \dots, h_n\}$
- **Objective:**
  - Find out a hypothesis  $h$  in  $H$  such that  $h(x)=c(x) \forall x$

# Hypothesis Space

- Set of all legal hypothesis defined by the chosen feature set and the chosen hypothesis language
- The space of all hypotheses that can, in principle, be output by a *learning algorithm*
- One way to think about a supervised learning machine is as a device that explores a “**hypothesis space**”
  - Each **setting of the parameters** in the machine is a **different hypothesis** about the function that **maps input vectors to output vectors**
- Given a set of data points, hypothesis  $h \in H$  is the **output of a learning algorithm**

# Hypothesis Space: Example

- $X = \{1, 2, 3, 4\}$ 
  - Set of all possible days
- **Features:**  $\langle \text{Sky, AirTemp, Humidity, Wind, Water} \rangle$
- $h_1$  : AirTemp = “cold” and Humidity = “high”
- $h_2$  : Sky = “sunny” and Water = “cool”

# Inductive Learning

- **Inductive learning:** Inducing a general function from training examples
  - Construct hypothesis  $h$  to agree with  $c$  on the training examples
  - A hypothesis is consistent if it agrees with all training examples
  - A hypothesis said to generalize well if it correctly predicts the value of  $y$  for novel example
- Inductive Learning is an ill Posed Problem:
  - Unless we see all possible examples the data is not sufficient for an inductive learning algorithm to find a unique solution

# Inductive Learning Hypothesis

- Any hypothesis  $h$  found to approximate the target function  $c$  well over a sufficiently large set of training examples  $D$  will *also approximate the target function well over other unobserved examples*



# Learning Issues

- What are good hypothesis spaces?
- Algorithms that work with the hypothesis spaces
- How to optimize accuracy over future data points
- How can we have confidence in the result? (How much training data?)

# Linear Regression

09/04/2024

**Koustav Rudra**

# Dataset of living area and price of houses in a city

Living area (feet <sup>2</sup> )	Price (1000\$)
2104	400
1600	330
2400	369
1416	232
3000	540
⋮	⋮

This is a training set.

How can we learn to **predict the prices of houses of other sizes** in the city, as a function of their living area?

# Dataset of living area and price of houses in a city

Living area (feet <sup>2</sup> )	Price (1000\$s)
2104	400
1600	330
2400	369
1416	232
3000	540
⋮	⋮

Example of supervised learning problem

When the target variable we are trying to predict is continuous,  
**regression** problem

# Dataset of living area and price of houses in a city

Living area (feet <sup>2</sup> )	Price (1000\$)
2104	400
1600	330
2400	369
1416	232
3000	540
$\vdots$	$\vdots$

$m$  = number of **training examples**

$x$ 's = input variables / features

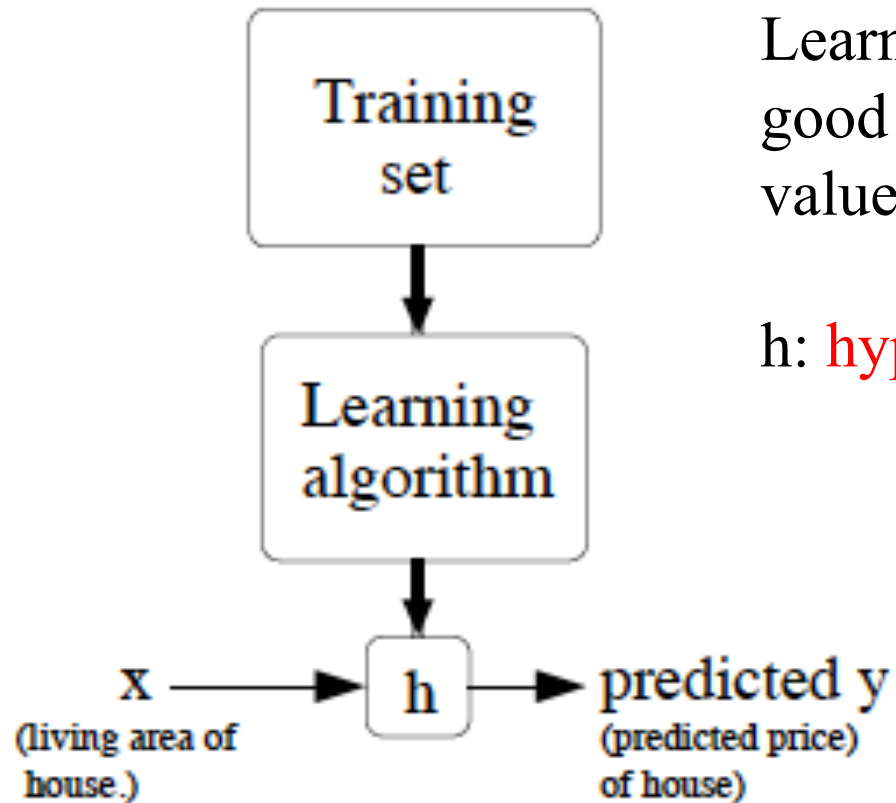
$y$ 's = output variables / "target" variables

$(x,y)$  - single training example

$(x^i, y^i)$  - specific example ( $i^{\text{th}}$  training example)

$i$  is an index to training set

# How to use the training set?



Learn a function  $h(x)$ , so that  $h(x)$  is a good predictor for the corresponding value of  $y$

$h$ : hypothesis function

# How to represent hypothesis $h$ ?

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

$\theta_i$  are **parameters**

- $\theta_0$  is zero condition

- $\theta_1$  is gradient

$\theta$ : vector of all the parameters

We assume  $y$  is a linear function of  $x$

Univariate linear regression

How to learn the values of the parameters?

# Digression: Multivariate linear regression

Living area (feet <sup>2</sup> )	#bedrooms	Price (1000\$s)
2104	3	400
1600	3	330
2400	3	369
1416	2	232
3000	4	540
⋮	⋮	⋮

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2$$



# How to represent hypothesis $h$ ?

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

$\theta_i$  are **parameters**


- $\theta_0$  is zero condition
- $\theta_1$  is gradient

We assume  $y$  is a linear function of  $x$

Univariate linear regression


How to learn the values of the parameters  $\theta_i$ ?

# Intuition of hypothesis function




- We are attempting to fit a straight line to the data in the training set
- Values of the parameters decide the equation of the straight line
- Which is the best straight line to fit the data?

# Intuition of hypothesis function



- Which is the best straight line to fit the data?
- How to learn the values of the parameters  $\theta_i$ ?
- Choose the parameters such that the prediction is close to the actual y-value for the training examples

# How good is the prediction given by the straight line?



$(x, y)$ : a training example

$(x, h(x))$ : prediction of the model

$(h(x) - y)$ : prediction error for this particular training example

**Minimize** the prediction error  
across **all training examples**

# Cost function

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

- Measure of how close the predictions are to the actual y-values
- Average over all the m training instances
- **Squared error cost function**  $J(\theta)$
- Choose parameters  $\theta$  so that  $J(\theta)$  is minimized

Hypothesis:  $h_{\theta}(x) = \theta_0 + \theta_1 x$

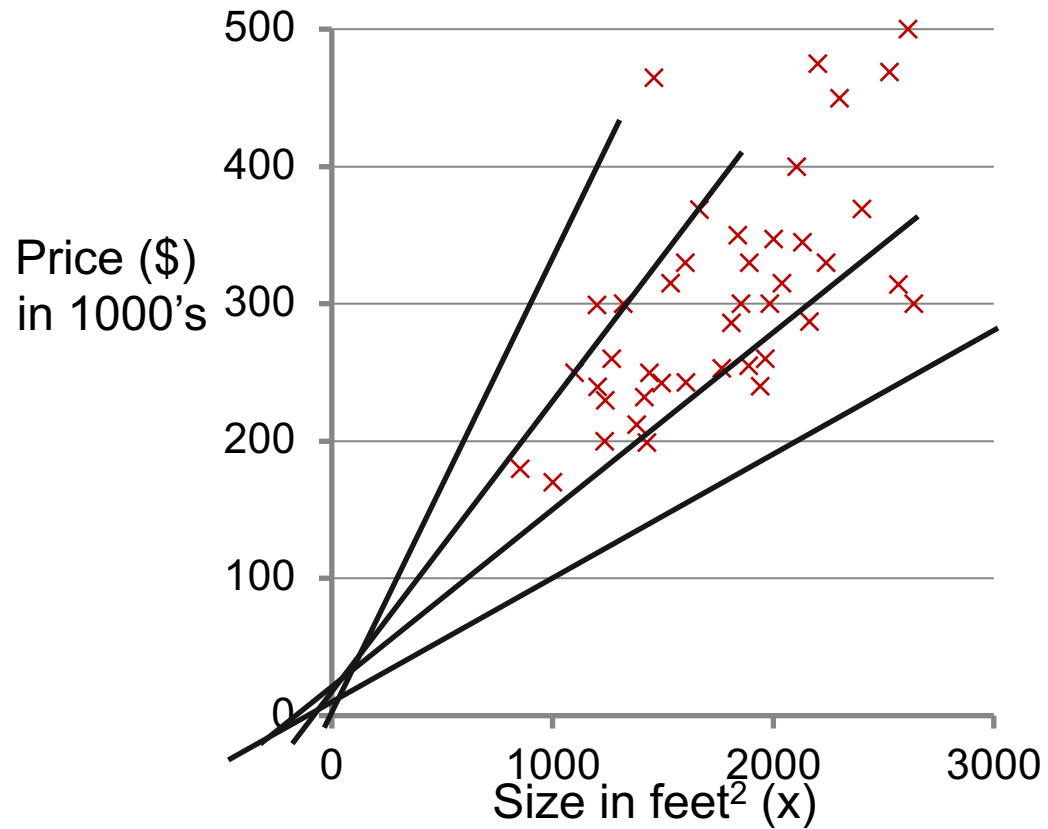
Parameters:  $\theta_0, \theta_1$

Cost Function:  $J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m \left( h_{\theta}(x^{(i)}) - y^{(i)} \right)^2$

Goal:  $\underset{\theta_0, \theta_1}{\text{minimize}} J(\theta_0, \theta_1)$


$$h_{\theta}(x)$$

(for fixed  $\theta_0, \theta_1$ , this is a function of  $x$ )



$$J(\theta_0, \theta_1)$$


(function of the parameters  $\theta_0, \theta_1$ )



For simplicity, assume  $\theta_0$  is a constant

# Contour plot or Contour figure

When  $J$  is a function of both  $\theta_0$  and  $\theta_1$






# Minimizing a function

- For now, let us consider some arbitrary function (not necessarily a cost function)
- Analytical minimization not scalable to complex functions of hundreds of parameters
- Algorithm called **gradient descent**
  - Efficient and scalable to thousands of parameters
  - Used in many applications of minimizing functions


Have some function  $J(\theta_0, \theta_1)$

Want  $\min_{\theta_0, \theta_1} J(\theta_0, \theta_1)$

- **Outline:**
  - Start with some  $\theta_0, \theta_1$
  - Keep changing  $\theta_0, \theta_1$  to reduce  $J(\theta_0, \theta_1)$  until we hopefully end up at a minimum
- Iterative method, similar to Newton-Raphson method for solving equations



If the function has multiple local minima, where one starts can decide which minimum is reached



# Gradient descent algorithm

repeat until convergence {

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1) \quad (\text{simultaneously update } j = 0 \text{ and } j = 1)$$

}

$\alpha$  is the **learning rate** – more on this later

# Gradient descent algorithm

repeat until convergence {  
     $\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$     (for  $j = 0$  and  $j = 1$ )  
}

---


Correct: Simultaneous update

temp0 :=  $\theta_0 - \alpha \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1)$   
temp1 :=  $\theta_1 - \alpha \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1)$   
 $\theta_0 :=$  temp0  
 $\theta_1 :=$  temp1

Incorrect:

temp0 :=  $\theta_0 - \alpha \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1)$   
 $\theta_0 :=$  temp0  
temp1 :=  $\theta_1 - \alpha \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1)$   
 $\theta_1 :=$  temp1

For simplicity, let us first consider a function of a single variable



$$\theta_1 := \theta_1 - \alpha \frac{\partial}{\partial \theta_1} J(\theta_1)$$

If the derivative is positive,  
reduce value of  $\theta_1$

If the derivative is negative,  
increase value of  $\theta_1$

# The learning rate

- Do we need to change learning rate over time?
  - No, Gradient descent can converge to a local minimum, even with the learning rate  $\alpha$  fixed
  - Step size adjusted automatically
- But, value needs to be chosen judiciously
  - If  $\alpha$  is too small, gradient descent can be slow to converge
  - If  $\alpha$  is too large, gradient descent can overshoot the minimum. It may fail to converge, or even diverge.



# Gradient descent for univariate linear regression

Gradient descent algorithm

repeat until convergence {  
     $\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$   
    (for  $j = 1$  and  $j = 0$ )  
}

Linear Regression Model

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m \left( h_{\theta}(x^{(i)}) - y^{(i)} \right)^2$$

# Gradient descent for univariate linear regression


$$\begin{array}{l} \text{repeat until convergence } \{ \\ \quad \left. \begin{array}{l} \theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) \\ \theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x^{(i)} \end{array} \right\} \begin{array}{l} \text{update} \\ \theta_0 \text{ and } \theta_1 \\ \text{simultaneously} \end{array} \\ \} \end{array}$$

# “Batch” Gradient Descent

- **“Batch”**: Each step of gradient descent uses all the training examples
  - At each estimate gradient on a batch of  $m$  samples
- There are other variations like “stochastic gradient descent” (used in learning over huge datasets)

# What about multiple local minima?

- The cost function in linear regression is always a **convex function** – always has a single global minimum
- So, gradient descent will always converge



# Linear Regression for multiple variables

## Multiple features (variables).

Size (feet <sup>2</sup> )	Number of bedrooms	Number of floors	Age of home (years)	Price (\$1000)
2104	5	1	45	460
1416	3	2	40	232
1534	3	2	30	315
852	2	1	36	178
...	...	...	...	...

## Multiple features (variables).

Size (feet <sup>2</sup> )	Number of bedrooms	Number of floors	Age of home (years)	Price (\$1000)
2104	5	1	45	460
1416	3	2	40	232
1534	3	2	30	315
852	2	1	36	178
...	...	...	...	...

Notation:

$n$  = number of features.  $m$  = number of training examples

$x^{(i)}$  = input (features) of  $i^{th}$  training example.

$x_j^{(i)}$  = value of feature  $j$  in  $i^{th}$  training example.

Hypothesis:

For univariate linear regression:

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

For **multi-variate linear regression**:

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_n x_n$$

For convenience of notation, define  $x_0 = 1$



Hypothesis:  $h_{\theta}(x) = \theta^T x = \theta_0 x_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_n x_n$

Parameters:  $\theta_0, \theta_1, \dots, \theta_n$

Cost function:

$$J(\theta_0, \theta_1, \dots, \theta_n) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Gradient descent:

Repeat {

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \dots, \theta_n)$$

} (simultaneously update for every  $j = 0, \dots, n$  )

# Gradient Descent

Previously (n=1):

Repeat {

$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})$$

$$\theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})x^{(i)}$$

(simultaneously update  $\theta_0, \theta_1$  )

}

New algorithm ( $n \geq 1$ ) :

Repeat {

$$\theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})x_j^{(i)}$$

simultaneously update  $\theta_j$  for  
 $j = 0, \dots, n$

}

---

...

# Gradient Descent

Previously (n=1):

Repeat {

$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})$$

$$\theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})x^{(i)}$$

(simultaneously update  $\theta_0, \theta_1$  )

}

New algorithm ( $n \geq 1$ ) :

Repeat {

$$\theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})x_j^{(i)}$$

simultaneously update  $\theta_j$  for  
 $j = 0, \dots, n$

}

---

$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})x_0^{(i)}$$

$$\theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})x_1^{(i)}$$

$$\theta_2 := \theta_2 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})x_2^{(i)}$$

...

# Practical aspects of applying gradient descent

# Feature Scaling

Idea: Make sure features are on a similar scale.

E.g.  $x_1$  = size (0-2000 feet<sup>2</sup>)

$x_2$  = number of bedrooms (1-5)

**Normalization wrt the maximum value:**

$$x_1 = \frac{\text{size (feet}^2\text{)}}{2000}$$

$$x_2 = \frac{\text{number of bedrooms}}{5}$$

# Feature Scaling

Idea: Make sure features are on a similar scale.

E.g.  $x_1$  = size (0-2000 feet<sup>2</sup>)

$x_2$  = number of bedrooms (1-5)

## Mean normalization:

Replace  $x_i$  with  $x_i - \mu_i$  to make features have approximately zero mean (Do not apply to  $x_0 = 1$  ).

## Other types of normalization:

$$x_1 = \frac{size - 1000}{2000}$$

$$x_2 = \frac{\#bedrooms - 2}{5}$$

$$-0.5 \leq x_1 \leq 0.5, -0.5 \leq x_2 \leq 0.5$$

# Is gradient descent working properly?

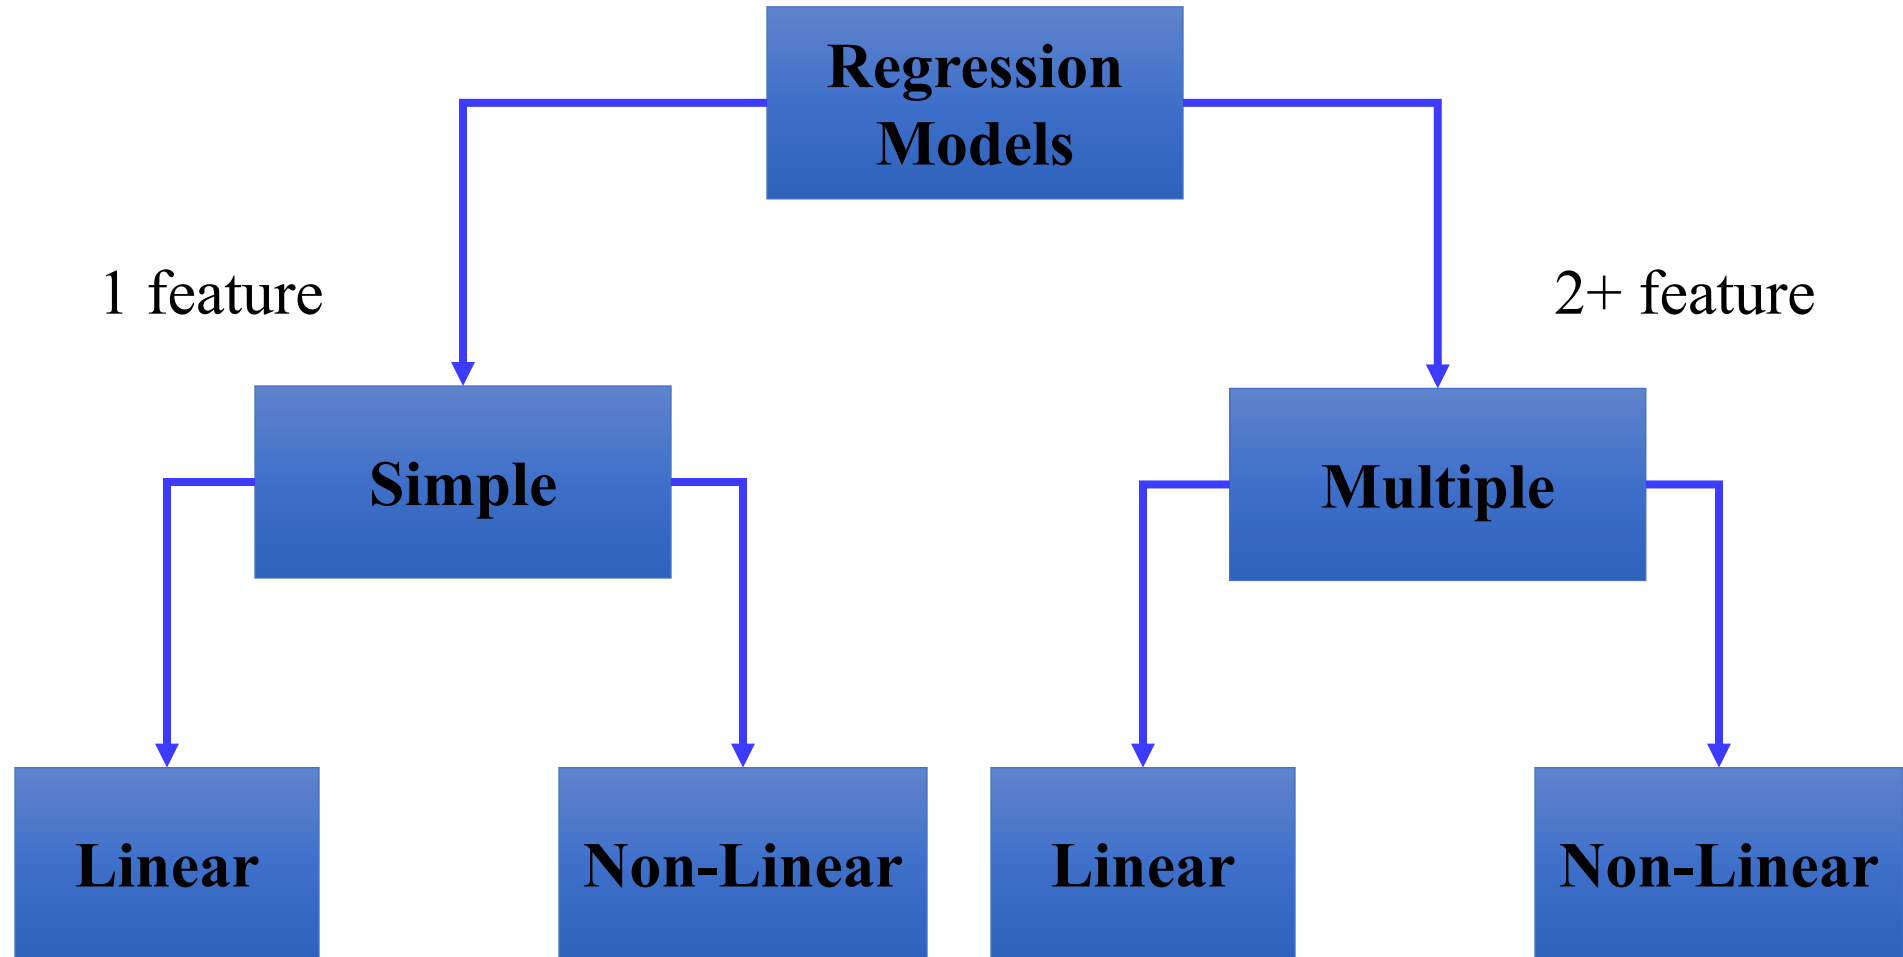
- Plot how  $J(\theta)$  changes with every iteration of gradient descent
- For sufficiently small learning rate,  $J(\theta)$  should decrease with every iteration
- If not, learning rate needs to be reduced
- However, too small learning rate means slow convergence

# When to end gradient descent?

- Example convergence test:
- Declare convergence if  $J(\theta)$  decreases by less than 0.001 in an iteration (assuming  $J(\theta)$  is decreasing in every iteration)



# Types of Regression Models



Thank You