

Decision Trees

15/04/2024

Koustav Rudra


Illustrating Classification Task

| Tid | Attrib1 | Attrib2 | Attrib3 | Class |
|-----|---------|---------|---------|-------|
| 1 | Yes | Large | 125k | No |
| 2 | No | Medium | 100k | No |
| 3 | No | Small | 70k | No |
| 4 | Yes | Medium | 120k | No |
| 5 | No | Large | 95k | Yes |
| 6 | No | Medium | 60k | No |
| 7 | Yes | Large | 220k | No |
| 8 | No | Small | 85k | Yes |
| 9 | No | Medium | 75k | No |
| 10 | No | Small | 90k | Yes |

Training Set

| Tid | Attrib1 | Attrib2 | Attrib3 | Class |
|-----|---------|---------|---------|-------|
| 1 | No | Small | 55k | ? |
| 2 | Yes | Medium | 80k | ? |
| 3 | Yes | Large | 110k | ? |
| 4 | No | Small | 95k | ? |
| 5 | No | Large | 67k | ? |

Test Set

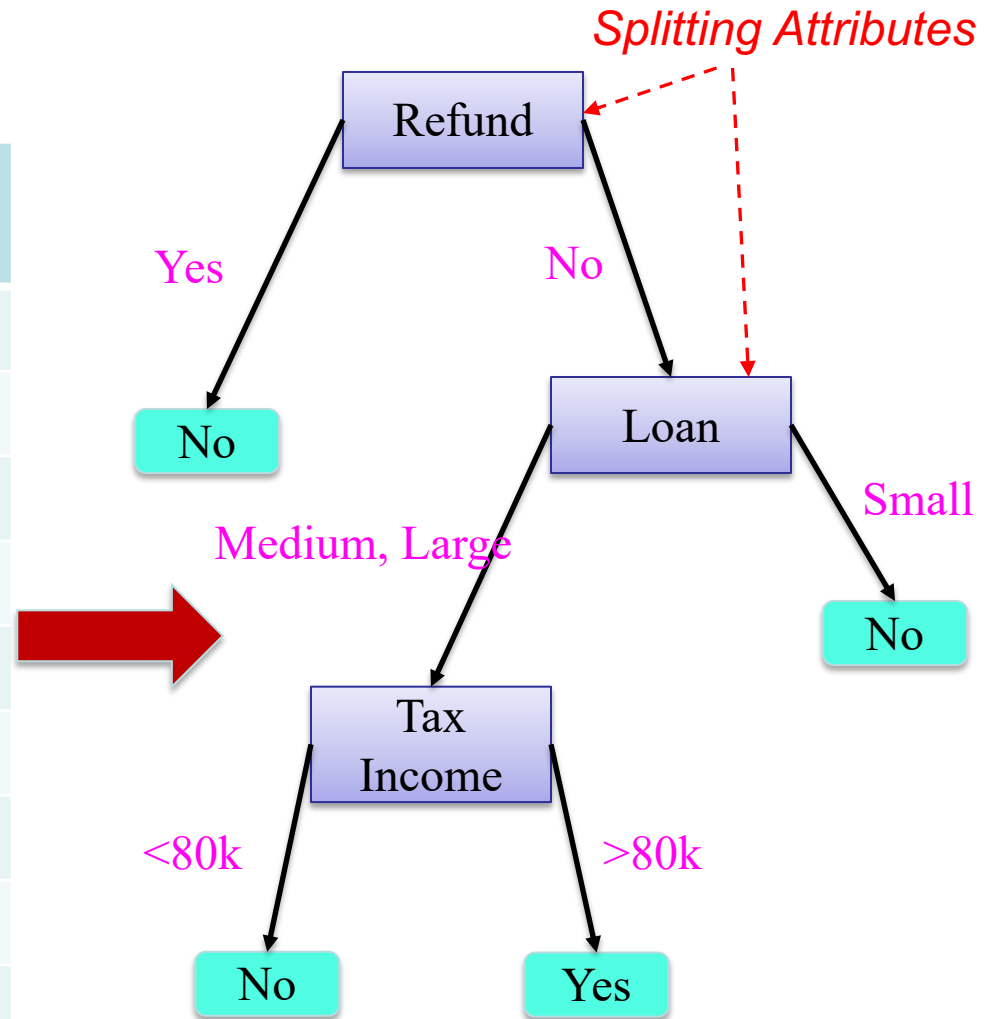


Intuition behind a decision tree

- Ask a series of questions about a given record
 - Each question is about one of the attributes
 - Answer to one question decides what question to ask next
(or if a next question is needed)
 - Continue asking questions until we can infer the class of the given record

Example of a Decision Tree

| Tid | categorical | | categorical | | class |
|-----|-------------|-------------|----------------|-------|-------|
| | Refund | Loan Status | Taxable Income | Cheat | |
| 1 | Yes | Medium | 125k | No | |
| 2 | No | Small | 100k | No | |
| 3 | No | Medium | 70k | No | |
| 4 | Yes | Small | 120k | No | |
| 5 | No | Large | 95k | Yes | |
| 6 | No | Small | 60k | No | |
| 7 | Yes | Large | 220k | No | |
| 8 | No | Medium | 85k | Yes | |
| 9 | No | Small | 75k | No | |
| 10 | No | Medium | 90k | Yes | |



Model: Decision Tree

Structure of a decision tree

- Decision tree: hierarchical structure
 - One **root node**: no incoming edge, zero or more outgoing edges
 - **Internal nodes**: exactly one incoming edge, two or more outgoing edges
 - **Leaf or terminal nodes**: exactly one incoming edge, no outgoing edge
- Each **leaf node** assigned a **class label**
- Each **non-leaf** node **contains a test condition** on one of the attributes

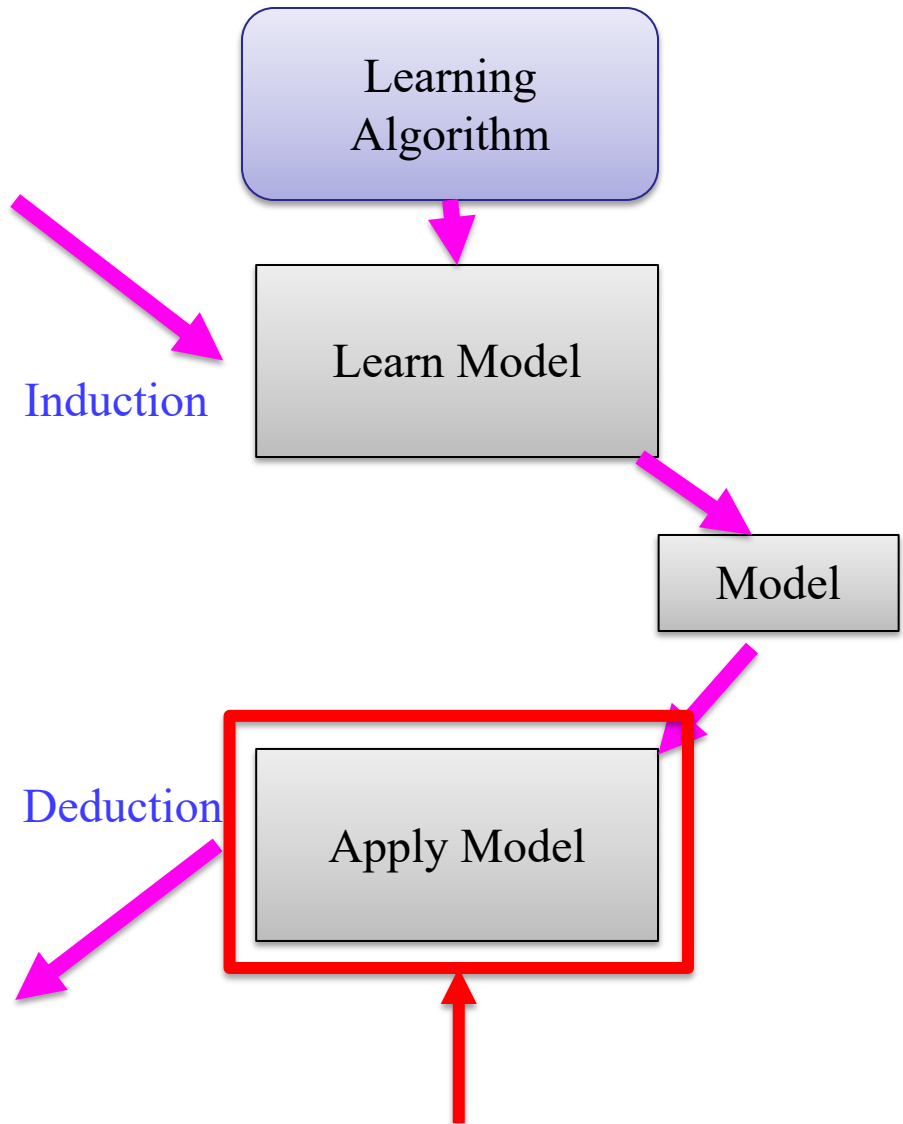
Applying a Decision-Tree Classifier

| Tid | Attrib1 | Attrib2 | Attrib3 | Class |
|-----|---------|---------|---------|-------|
| 1 | Yes | Large | 125k | No |
| 2 | No | Medium | 100k | No |
| 3 | No | Small | 70k | No |
| 4 | Yes | Medium | 120k | No |
| 5 | No | Large | 95k | Yes |
| 6 | No | Medium | 60k | No |
| 7 | Yes | Large | 220k | No |
| 8 | No | Small | 85k | Yes |
| 9 | No | Medium | 75k | No |
| 10 | No | Small | 90k | Yes |

Training Set

| Tid | Attrib1 | Attrib2 | Attrib3 | Class |
|-----|---------|---------|---------|-------|
| 1 | No | Small | 55k | ? |
| 2 | Yes | Medium | 80k | ? |
| 3 | Yes | Large | 110k | ? |
| 4 | No | Small | 95k | ? |
| 5 | No | Large | 67k | ? |

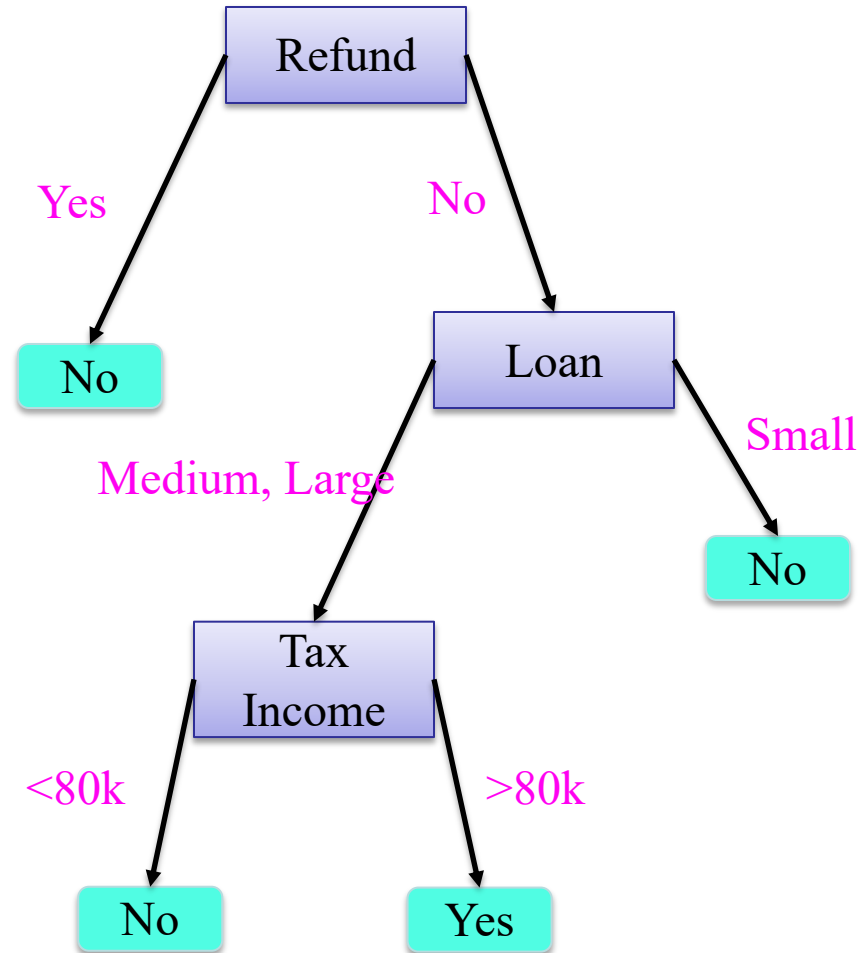
Test Set



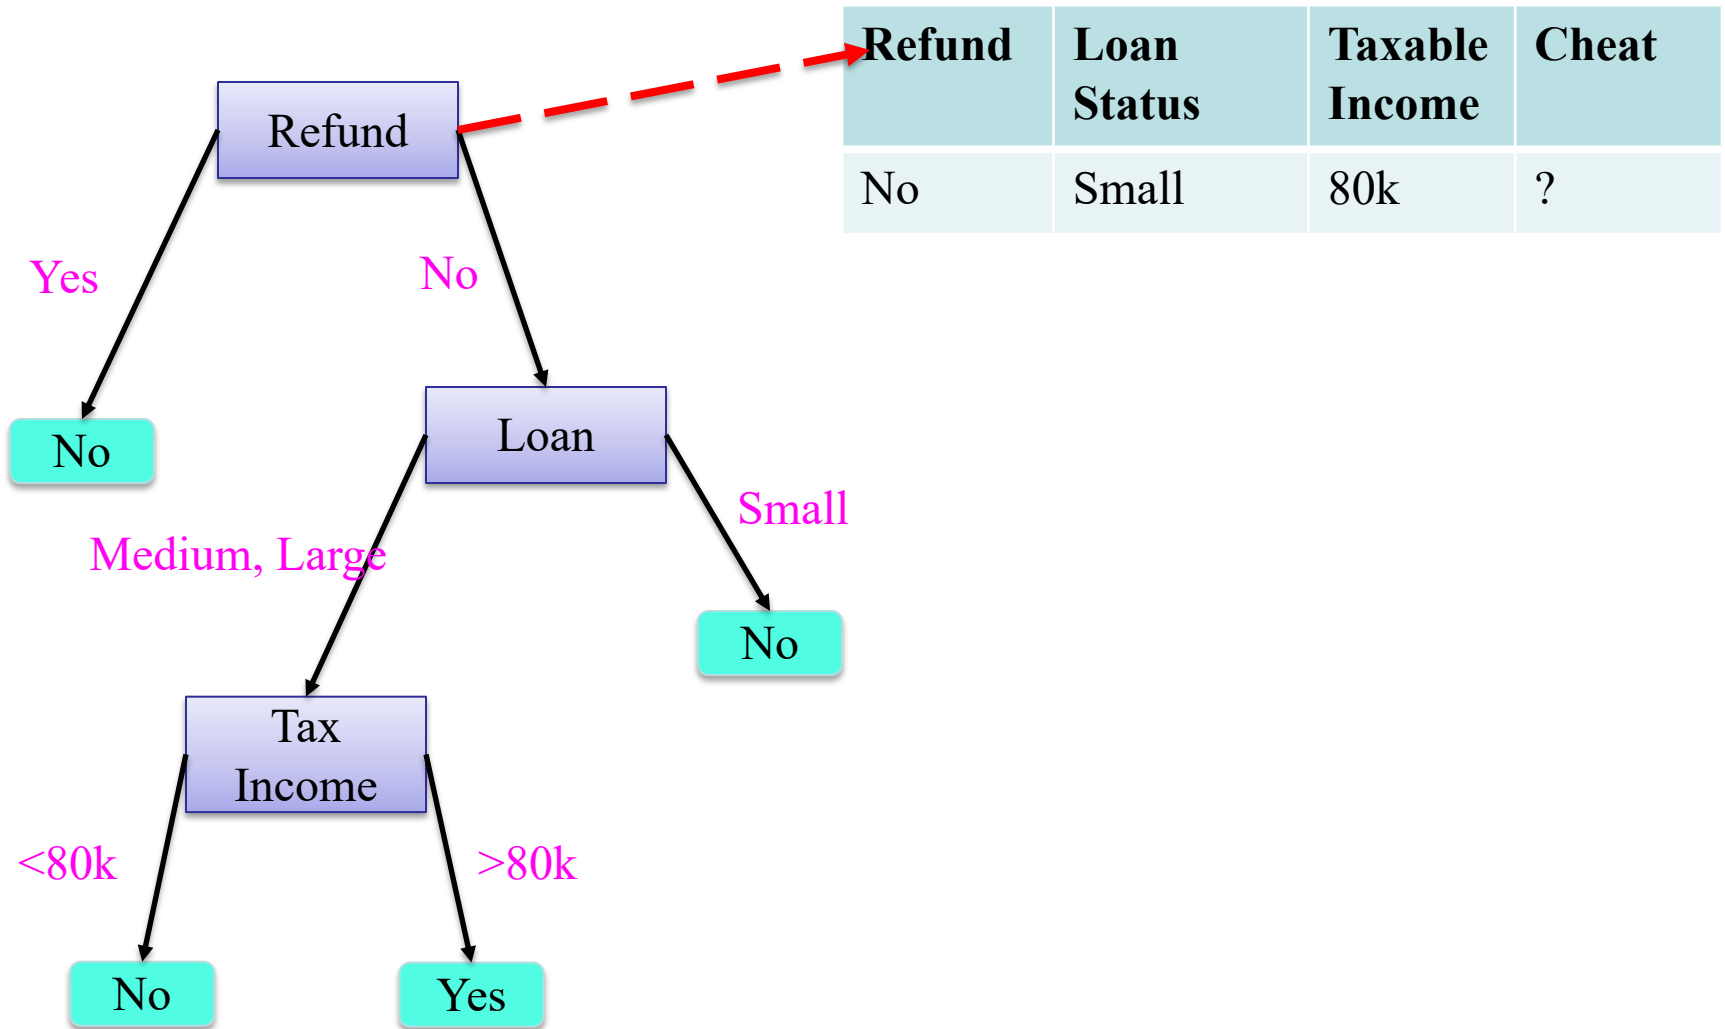
Applying Model to Test Data

| Refund | Loan Status | Taxable Income | Cheat |
|--------|-------------|----------------|-------|
| No | Small | 80k | ? |

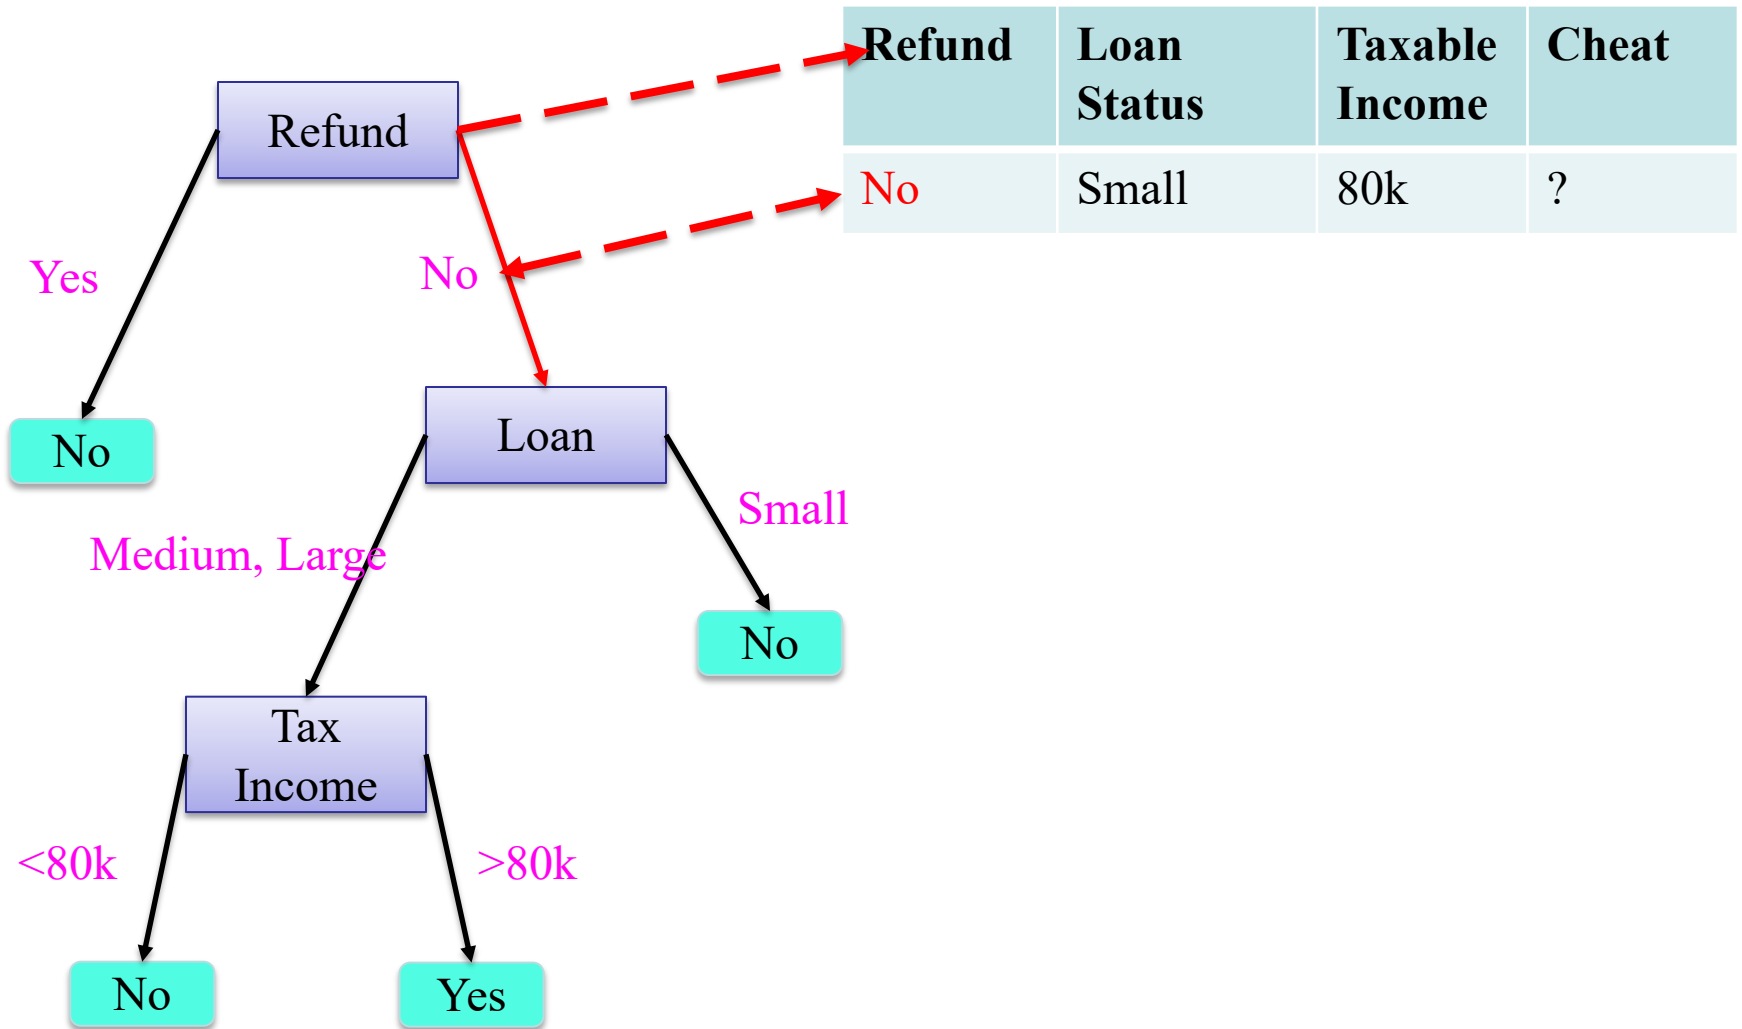
Once a decision tree has been constructed (learned), it is easy to apply it to test data



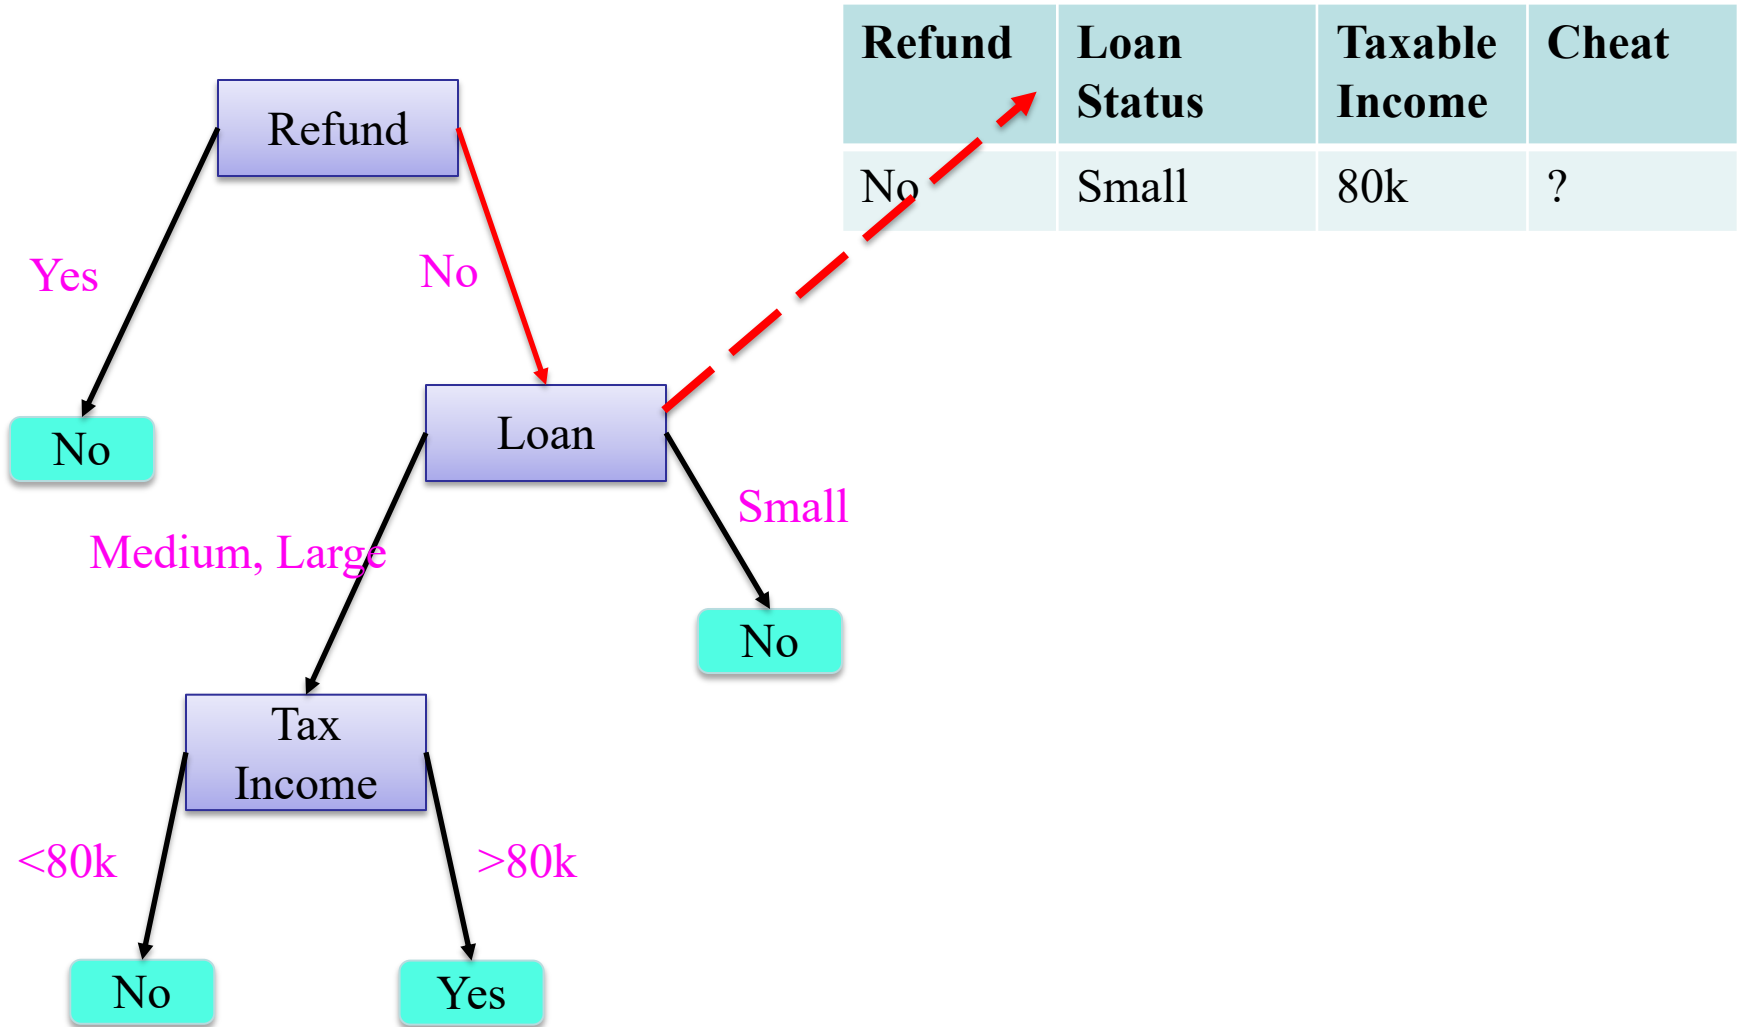
Applying Model to Test Data



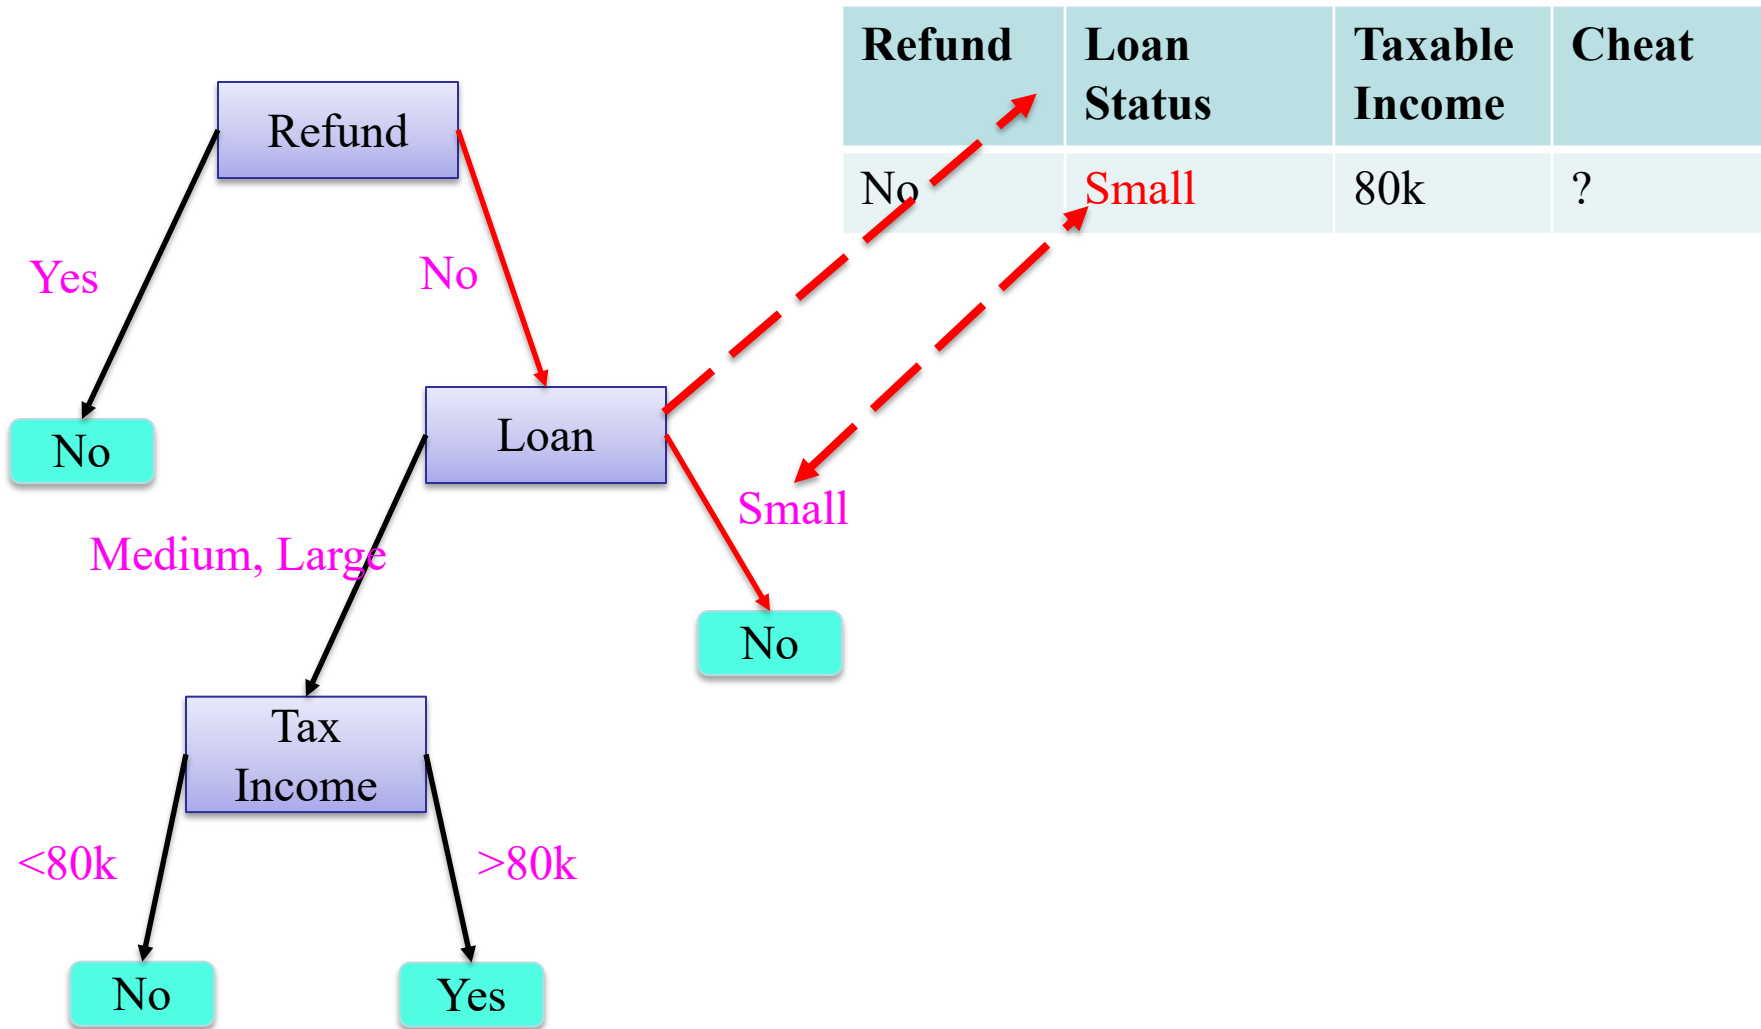
Applying Model to Test Data



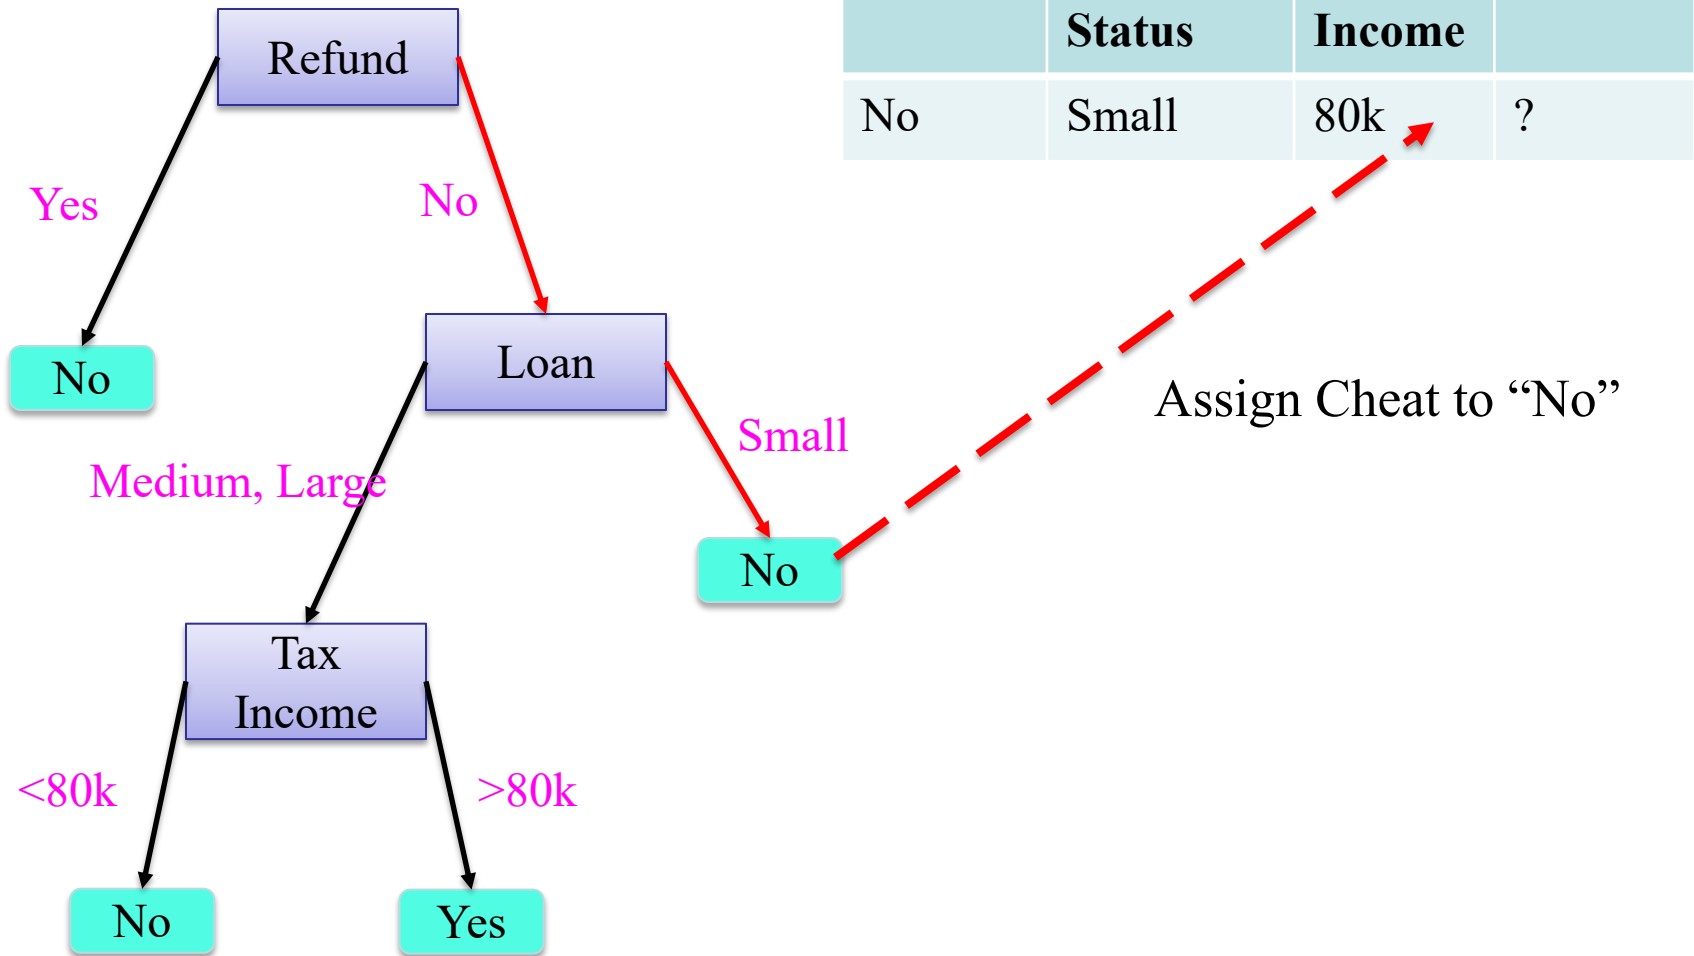
Applying Model to Test Data



Applying Model to Test Data



Applying Model to Test Data



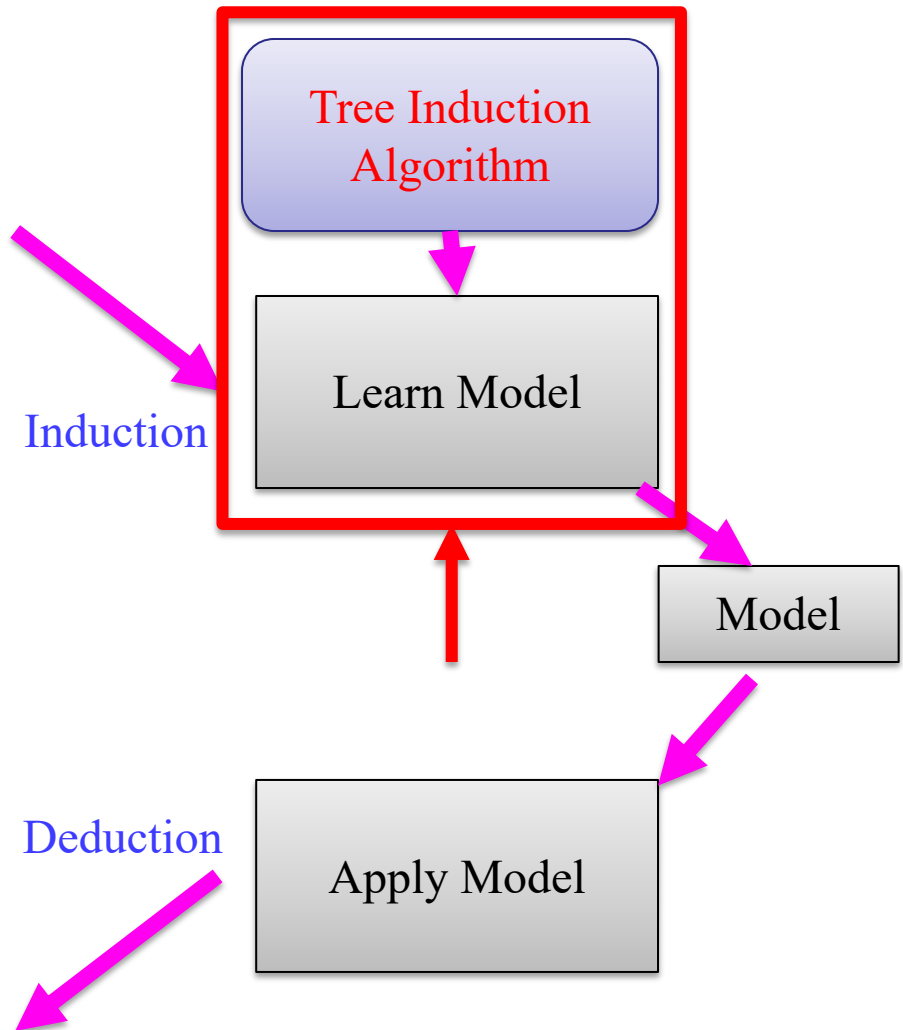
Learning a Decision-Tree Classifier

| Tid | Attrib1 | Attrib2 | Attrib3 | Class |
|-----|---------|---------|---------|-------|
| 1 | Yes | Large | 125k | No |
| 2 | No | Medium | 100k | No |
| 3 | No | Small | 70k | No |
| 4 | Yes | Medium | 120k | No |
| 5 | No | Large | 95k | Yes |
| 6 | No | Medium | 60k | No |
| 7 | Yes | Large | 220k | No |
| 8 | No | Small | 85k | Yes |
| 9 | No | Medium | 75k | No |
| 10 | No | Small | 90k | Yes |

Training Set

| Tid | Attrib1 | Attrib2 | Attrib3 | Class |
|-----|---------|---------|---------|-------|
| 1 | No | Small | 55k | ? |
| 2 | Yes | Medium | 80k | ? |
| 3 | Yes | Large | 110k | ? |
| 4 | No | Small | 95k | ? |
| 5 | No | Large | 67k | ? |

Test Set

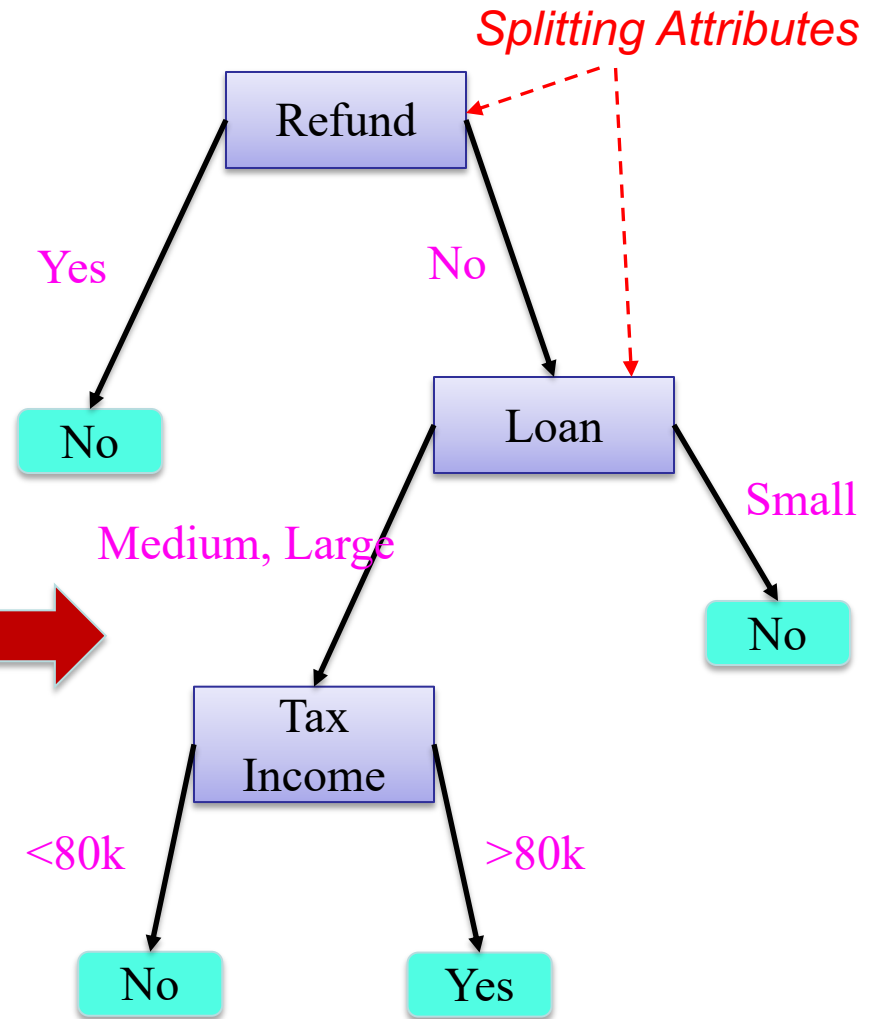


How to learn a decision tree?

A Decision Tree

| Tid | Refund | Loan Status | Taxable Income | Cheat |
|-----|--------|-------------|----------------|-------|
| 1 | Yes | Medium | 125k | No |
| 2 | No | Small | 100k | No |
| 3 | No | Medium | 70k | No |
| 4 | Yes | Small | 120k | No |
| 5 | No | Large | 95k | Yes |
| 6 | No | Small | 60k | No |
| 7 | Yes | Large | 220k | No |
| 8 | No | Medium | 85k | Yes |
| 9 | No | Small | 75k | No |
| 10 | No | Medium | 90k | Yes |

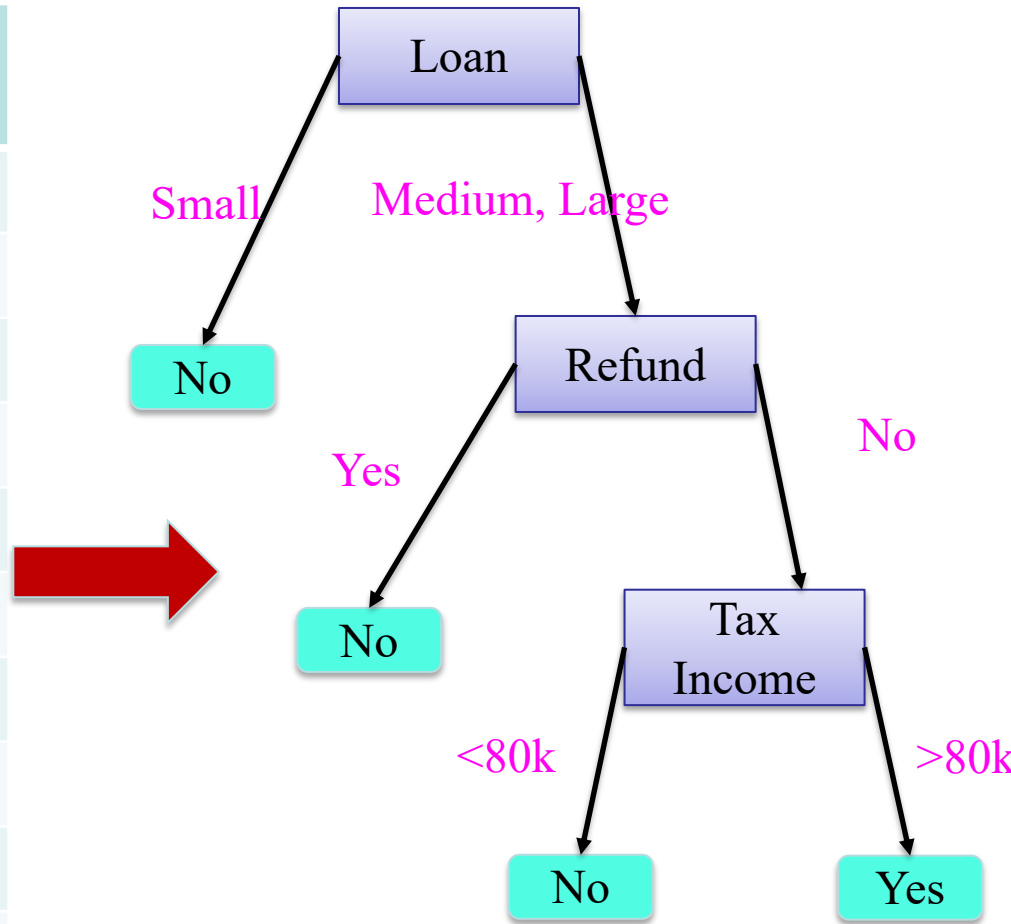
Training data



Model: Decision Tree

Another Decision Tree on same dataset

| Tid | Refund | Loan Status | Taxable Income | Cheat |
|-----|--------|-------------|----------------|-------|
| 1 | Yes | Medium | 125k | No |
| 2 | No | Small | 100k | No |
| 3 | No | Medium | 70k | No |
| 4 | Yes | Small | 120k | No |
| 5 | No | Large | 95k | Yes |
| 6 | No | Small | 60k | No |
| 7 | Yes | Large | 220k | No |
| 8 | No | Medium | 85k | Yes |
| 9 | No | Small | 75k | No |
| 10 | No | Medium | 90k | Yes |



There could be more than one tree that fits the same data!


Decision Tree Induction

- Many Algorithms:
 - Hunt's Algorithm (one of the earliest)
 - CART
 - ID3, C4.5
 - SLIQ, SPRINT

General Structure of Hunt's Algorithm

- Let D_t be the set of training records that reach a node t
- **General Procedure:**
 - If D_t contains records that all belong the same class y_t , then t is a leaf node labeled as y_t
 - If D_t is an empty set, then t is a leaf node labeled by the default class y_d
 - If D_t contains records that belong to more than one class, *use an attribute test to split the data into smaller subsets*. Recursively apply the procedure to each subset

| Tid | Refund | Loan Status | Taxable Income | Cheat |
|-----|--------|-------------|----------------|-------|
| 1 | Yes | Medium | 125k | No |
| 2 | No | Small | 100k | No |
| 3 | No | Medium | 70k | No |
| 4 | Yes | Small | 120k | No |
| 5 | No | Large | 95k | Yes |
| 6 | No | Small | 60k | No |
| 7 | Yes | Large | 220k | No |
| 8 | No | Medium | 85k | Yes |
| 9 | No | Small | 75k | No |
| 10 | No | Medium | 90k | Yes |



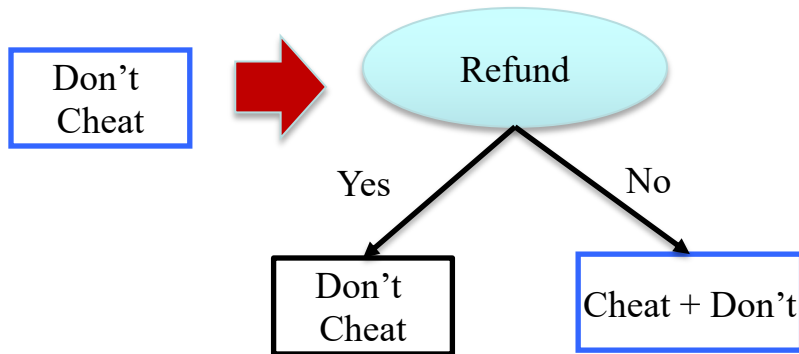
Hunt's Algorithm

Don't
Cheat + Cheat

| Tid | Refund | Loan Status | Taxable Income | Cheat |
|-----|--------|-------------|----------------|-------|
| 1 | Yes | Medium | 125k | No |
| 2 | No | Small | 100k | No |
| 3 | No | Medium | 70k | No |
| 4 | Yes | Small | 120k | No |
| 5 | No | Large | 95k | Yes |
| 6 | No | Small | 60k | No |
| 7 | Yes | Large | 220k | No |
| 8 | No | Medium | 85k | Yes |
| 9 | No | Small | 75k | No |
| 10 | No | Medium | 90k | Yes |

Default class is “Don't cheat”
since it is the majority class in
the dataset

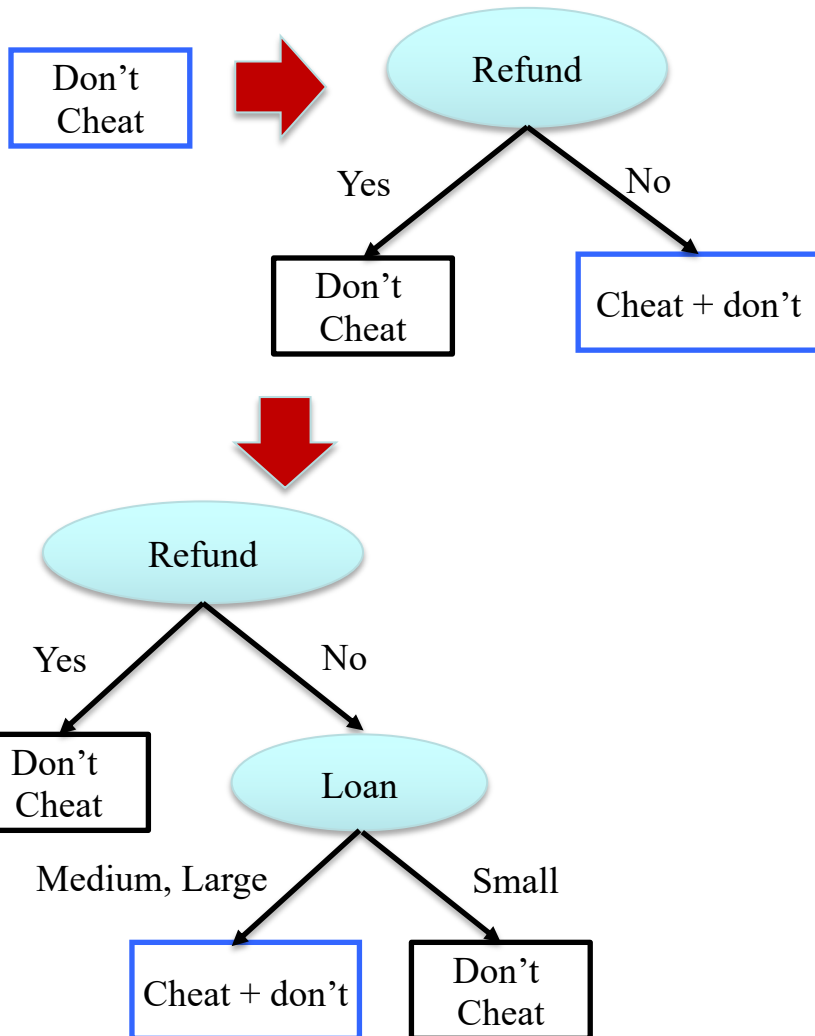
Hunt's Algorithm



| Tid | Refund | Loan Status | Taxable Income | Cheat |
|-----|--------|-------------|----------------|-------|
| 1 | Yes | Medium | 125k | No |
| 2 | No | Small | 100k | No |
| 3 | No | Medium | 70k | No |
| 4 | Yes | Small | 120k | No |
| 5 | No | Large | 95k | Yes |
| 6 | No | Small | 60k | No |
| 7 | Yes | Large | 220k | No |
| 8 | No | Medium | 85k | Yes |
| 9 | No | Small | 75k | No |
| 10 | No | Medium | 90k | Yes |

For now, assume that “**Refund**” has been decided to be **the best attribute** for **splitting in some way** (to be discussed soon)

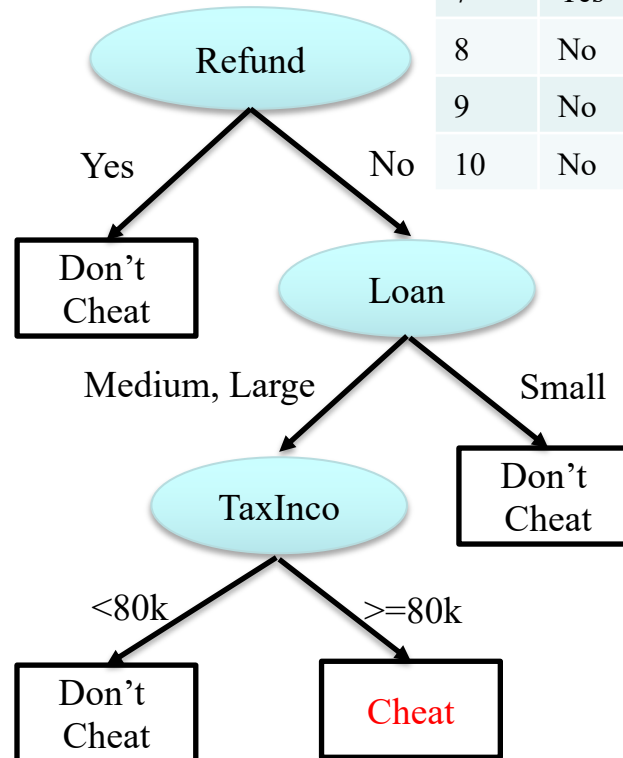
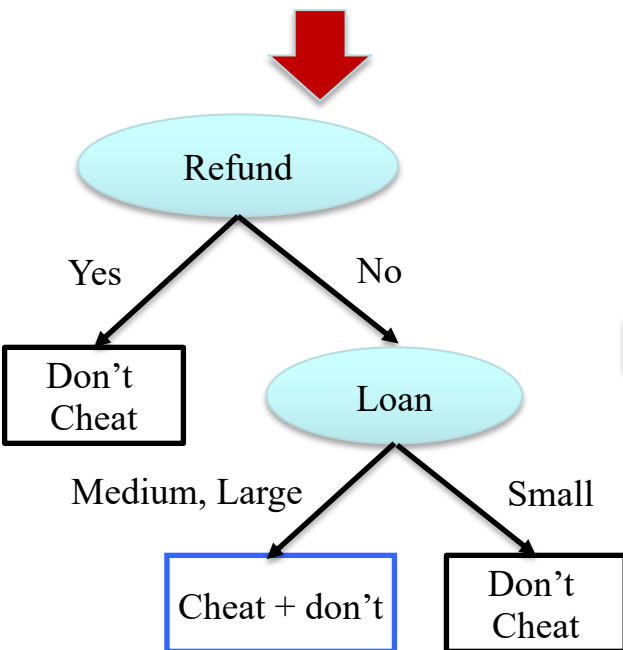
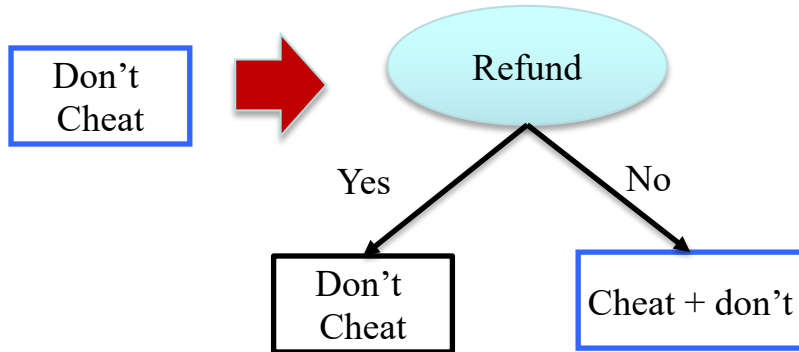
Hunt's Algorithm



| Tid | Refund | Loan Status | Taxable Income | Cheat |
|-----|--------|-------------|----------------|-------|
| 1 | Yes | Medium | 125k | No |
| 2 | No | Small | 100k | No |
| 3 | No | Medium | 70k | No |
| 4 | Yes | Small | 120k | No |
| 5 | No | Large | 95k | Yes |
| 6 | No | Small | 60k | No |
| 7 | Yes | Large | 220k | No |
| 8 | No | Medium | 85k | Yes |
| 9 | No | Small | 75k | No |
| 10 | No | Medium | 90k | Yes |

Hunt's Algorithm

| Tid | Refund | Loan Status | Taxable Income | Cheat |
|-----|--------|-------------|----------------|-------|
| 1 | Yes | Medium | 125k | No |
| 2 | No | Small | 100k | No |
| 3 | No | Medium | 70k | No |
| 4 | Yes | Small | 120k | No |
| 5 | No | Large | 95k | Yes |
| 6 | No | Small | 60k | No |
| 7 | Yes | Large | 220k | No |
| 8 | No | Medium | 85k | Yes |
| 9 | No | Small | 75k | No |
| 10 | No | Medium | 90k | Yes |




How to Specify Test Condition?


- Depends on **attribute types**
 - **Nominal**: two or more distinct values (special case: binary)
E.g., Loan status: {small, medium, large}
 - **Ordinal**: two or more distinct values that have an ordering.
E.g. shirt size: {S, M, L, XL}
 - **Continuous**: continuous range of values
- Depends on **number of ways to split**
 - 2-way split
 - Multi-way split

Splitting Based on Nominal Attributes


- **Multi-way split:** Use as many partitions as distinct values.



- **Binary split:** Divides values into two subsets.
Need to find optimal partitioning




OR




Splitting Based on Ordinal Attributes

- **Multi-way split:** Use as many partitions as distinct values.



- **Binary split:** Divides values into two subsets.
Need to find optimal partitioning.



- **What about this split?**

Tree Induction

- Greedy strategy
 - Split the records based on an attribute test that optimizes certain criterion
- Issues
 - Determine how to split the records
 - How to specify the attribute test condition?
 - How to determine the best split?
 - Determine when to stop splitting

Decision Trees

Finding Best Attribute

15/04/2024


Measures of Node Impurity

- Gini Index
- Entropy
- Misclassification error



How to Find the Best Split

Before Splitting:

| | | |
|----|------------|------|
| C0 | N00 | → M0 |
| C1 | N01 | |



| | | | |
|----|------------|----|------------|
| C0 | N10 | C0 | N20 |
| C1 | N11 | C1 | N21 |



| | | | |
|----|------------|----|------------|
| C0 | N30 | C0 | N40 |
| C1 | N31 | C1 | N41 |



$$\text{Gain} = M0 - M12 \text{ vs } M0 - M34$$

Measures of Node Impurity

- Gini Index
- Entropy
- Misclassification error

Alternative Splitting Criteria based on INFO

- Entropy at a given node t:

$$Entropy(t) = -\sum_j p(j | t) \log_2 p(j | t)$$

- $p(j | t)$ is the relative frequency of class j at node t
- Measures homogeneity of a node
- Entropy of sample S: Average optimal number of bits to encode information about certainty/uncertainty about S

Examples for computing Entropy

$$Entropy(t) = -\sum_j p(j | t) \log_2 p(j | t)$$

| | |
|----|----------|
| C1 | 0 |
| C2 | 6 |

$$P(C1) = 0/6 = 0 \quad P(C2) = 6/6 = 1$$

$$Entropy = -0 \log 0 - 1 \log 1 = -0 - 0 = 0$$

| | |
|----|----------|
| C1 | 1 |
| C2 | 5 |

$$P(C1) = 1/6 \quad P(C2) = 5/6$$

$$Entropy = - (1/6) \log_2 (1/6) - (5/6) \log_2 (1/6) = 0.65$$

| | |
|----|----------|
| C1 | 2 |
| C2 | 4 |

$$P(C1) = 2/6 \quad P(C2) = 4/6$$

$$Entropy = - (2/6) \log_2 (2/6) - (4/6) \log_2 (4/6) = 0.92$$

Alternative Splitting Criteria based on INFO

- Entropy at a given node t:

$$Entropy(t) = -\sum_j p(j | t) \log_2 p(j | t)$$


- $p(j | t)$ is the relative frequency of class j at node t
- Measures homogeneity of a node
 - Maximum ($\log n_c$) when records are equally distributed among all classes
 - implying least information
 - Minimum (0.0) when all records belong to one class,
 - implying most information

Information Gain

- Measures how well a given attribute separates the training examples according to their target classification
- This measure is used to select among the candidate attributes at each step while growing the tree
- Gain is measure of how much we can reduce uncertainty (Value lies between $[0,1]$)

Information Gain

S: [9+, 5-]
E: 0.940




$$\text{Entropy}(3+, 4-) = - (3/7)\log(3/7) - (4/7)\log(4/7) = 0.985$$

$$\text{Entropy}(6+, 1-) = - (6/7)\log(6/7) - (1/7)\log(1/7) = 0.592$$

$$\text{Gain}(S, A1) = 0.940 - (7/14)*0.985 - (7/14)*0.592 = 0.151$$

S: [9+, 5-]
E: 0.940



$$\text{Entropy}(6+, 2-) = - (6/8)\log(6/8) - (2/8)\log(2/8) = 0.811$$

$$\text{Entropy}(3+, 3-) = - (3/6)\log(3/6) - (3/6)\log(3/6) = 1.0$$

$$\text{Gain}(S, A2) = 0.940 - (8/14)*0.811 - (6/14)*1.0 = 0.048$$

Splitting Based on INFO...

- Information Gain:

$$GAIN_{split} = Entropy(p) - \left(\sum_{i=1}^k \frac{n_i}{n} Entropy(i) \right)$$

- Parent Node p is split into k partitions;
- n_i is number of records in partition i
- Measures Reduction in Entropy achieved because of the split
Choose the split that achieves most reduction (maximizes GAIN)
- Used in ID3 and C4.5
- **Disadvantage:** Tends to prefer splits that result in large number of partitions, each being small but pure

Splitting Based on INFO...

- Gain Ratio:

$$GainRATIO_{split} = \frac{GAIN_{Split}}{SplitINFO} \quad SplitINFO = -\sum_{i=1}^k \frac{n_i}{n} \log \frac{n_i}{n}$$

- Parent Node, p is split into k partitions
- n_i is the number of records in partition i
- Adjusts Information Gain by the entropy of the partitioning (SplitINFO)
 - Higher entropy partitioning (large number of small partitions) is penalized!
- Used in C4.5
- Designed to overcome the disadvantage of Information Gain

Stopping Criteria for Tree Induction

- Stop expanding a node when all the records belong to the same class
- Stop expanding a node when all the records have similar attribute values (if different class values, then usually assign the majority class)
- Early termination, usually to prevent overfitting (to be discussed later)

ID3

ID3(Examples, Target_attribute, Attributes)

- Create a Root node for the tree
- If all examples are positive, Returns single-node tree Root with label +
- If all examples are negative, Returns single-node tree Root with label –
- If Attributes is empty, Returns single-node tree Root, with label = most common value of Target_attribute in Examples

ID3

ID3(Examples, Target_attribute, Attributes)

- Begin
 - $A \leftarrow$ Best attribute from Examples
 - The decision attribute for Root $\leftarrow A$
 - For each possible value, v_i , of A ,
 - Add a new branch below Root, corresponding to $A = v_i$
 - Examples_ v_i subset of examples with $v_i = A$
 - If Examples_ v_i is empty
 - Add a leaf node with label = most common value of Target_attribute in Examples
 - Else below this new branch add the subtree
 - ID3(Examples_ v_i , Target_attribute, Attributes- $\{A\}$)
- End
- Return Root

Thank You